

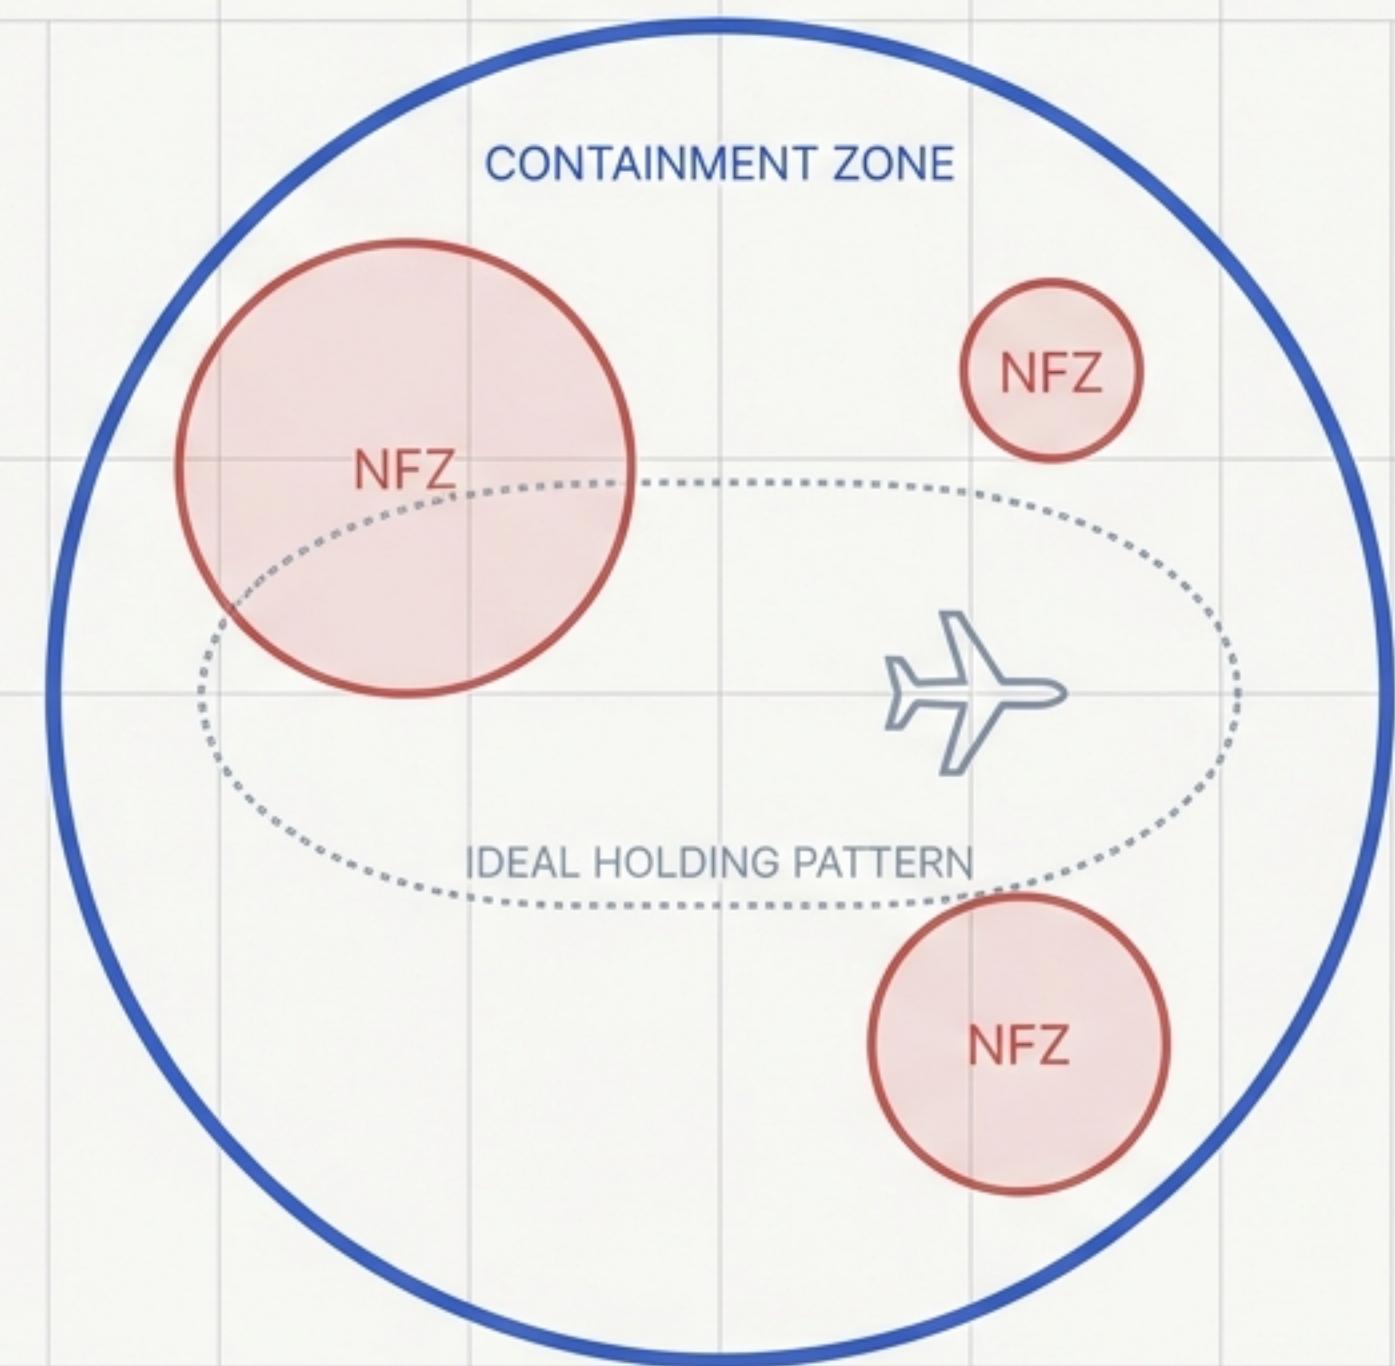


# **HoldingAgentLidarEnv: A Blueprint for Portable Reinforcement Learning**

Designing simulator-agnostic environments with a clean,  
testable architecture.

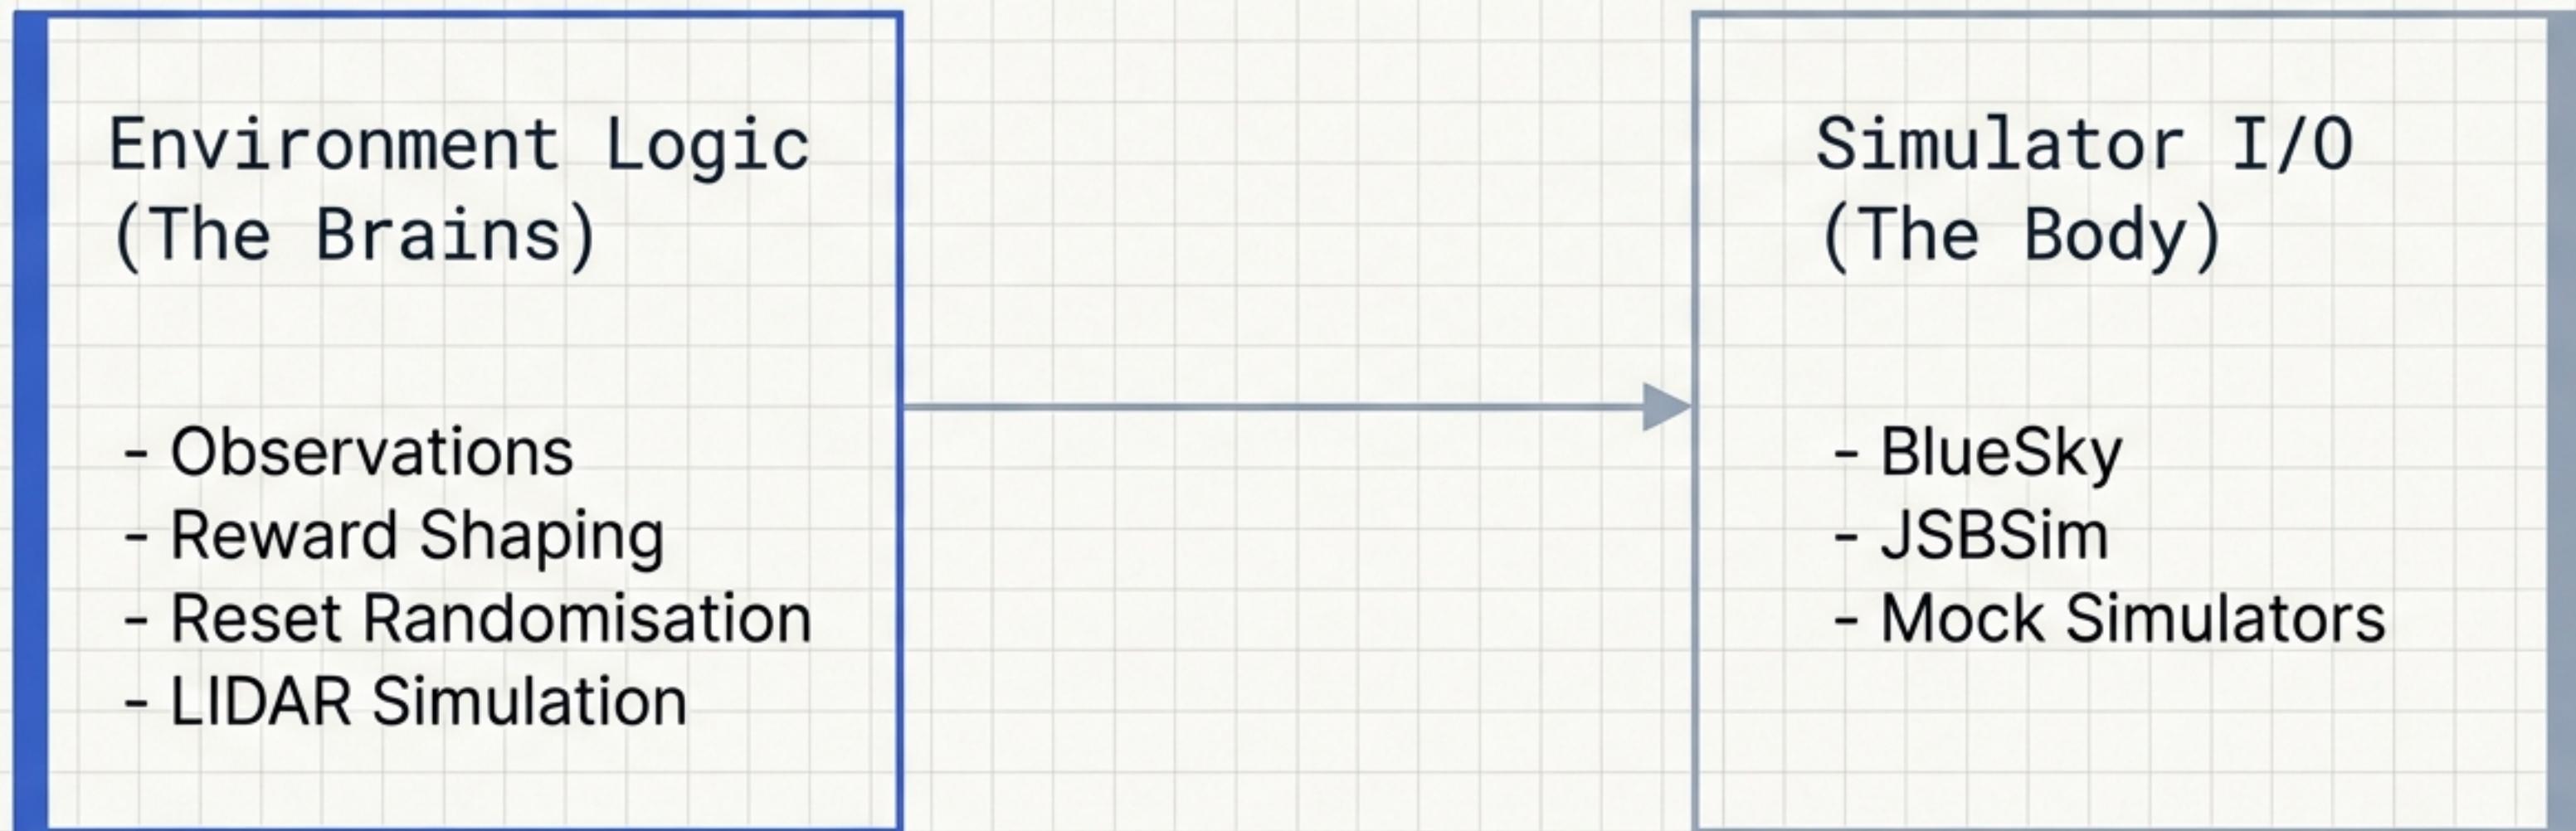
# THE MISSION: TRAIN AN AGENT FOR A COMPLEX AEROSPACE TASK

To train an aircraft agent to hold a pattern inside a circular containment zone while avoiding circular No-Fly Zones (NFZs) within a robust, simulator-independent framework.



# THE CORE PHILOSOPHY IS A CLEAN SEPARATION OF CONCERNS

The environment's core logic is completely independent of any specific simulator.  
The design intentionally separates simulator I/O from the environment's logic.



# THIS ARCHITECTURAL DISCIPLINE DELIVERS STRATEGIC ADVANTAGES



## TESTABILITY

Mock the simulator for unit tests. Test environment logic without needing to run a full, heavy simulator.



## PORTABILITY

Swap simulator backends with minimal code changes. Move from BlueSky to JSBSim or another compliant simulator easily.



## MAINTAINABILITY

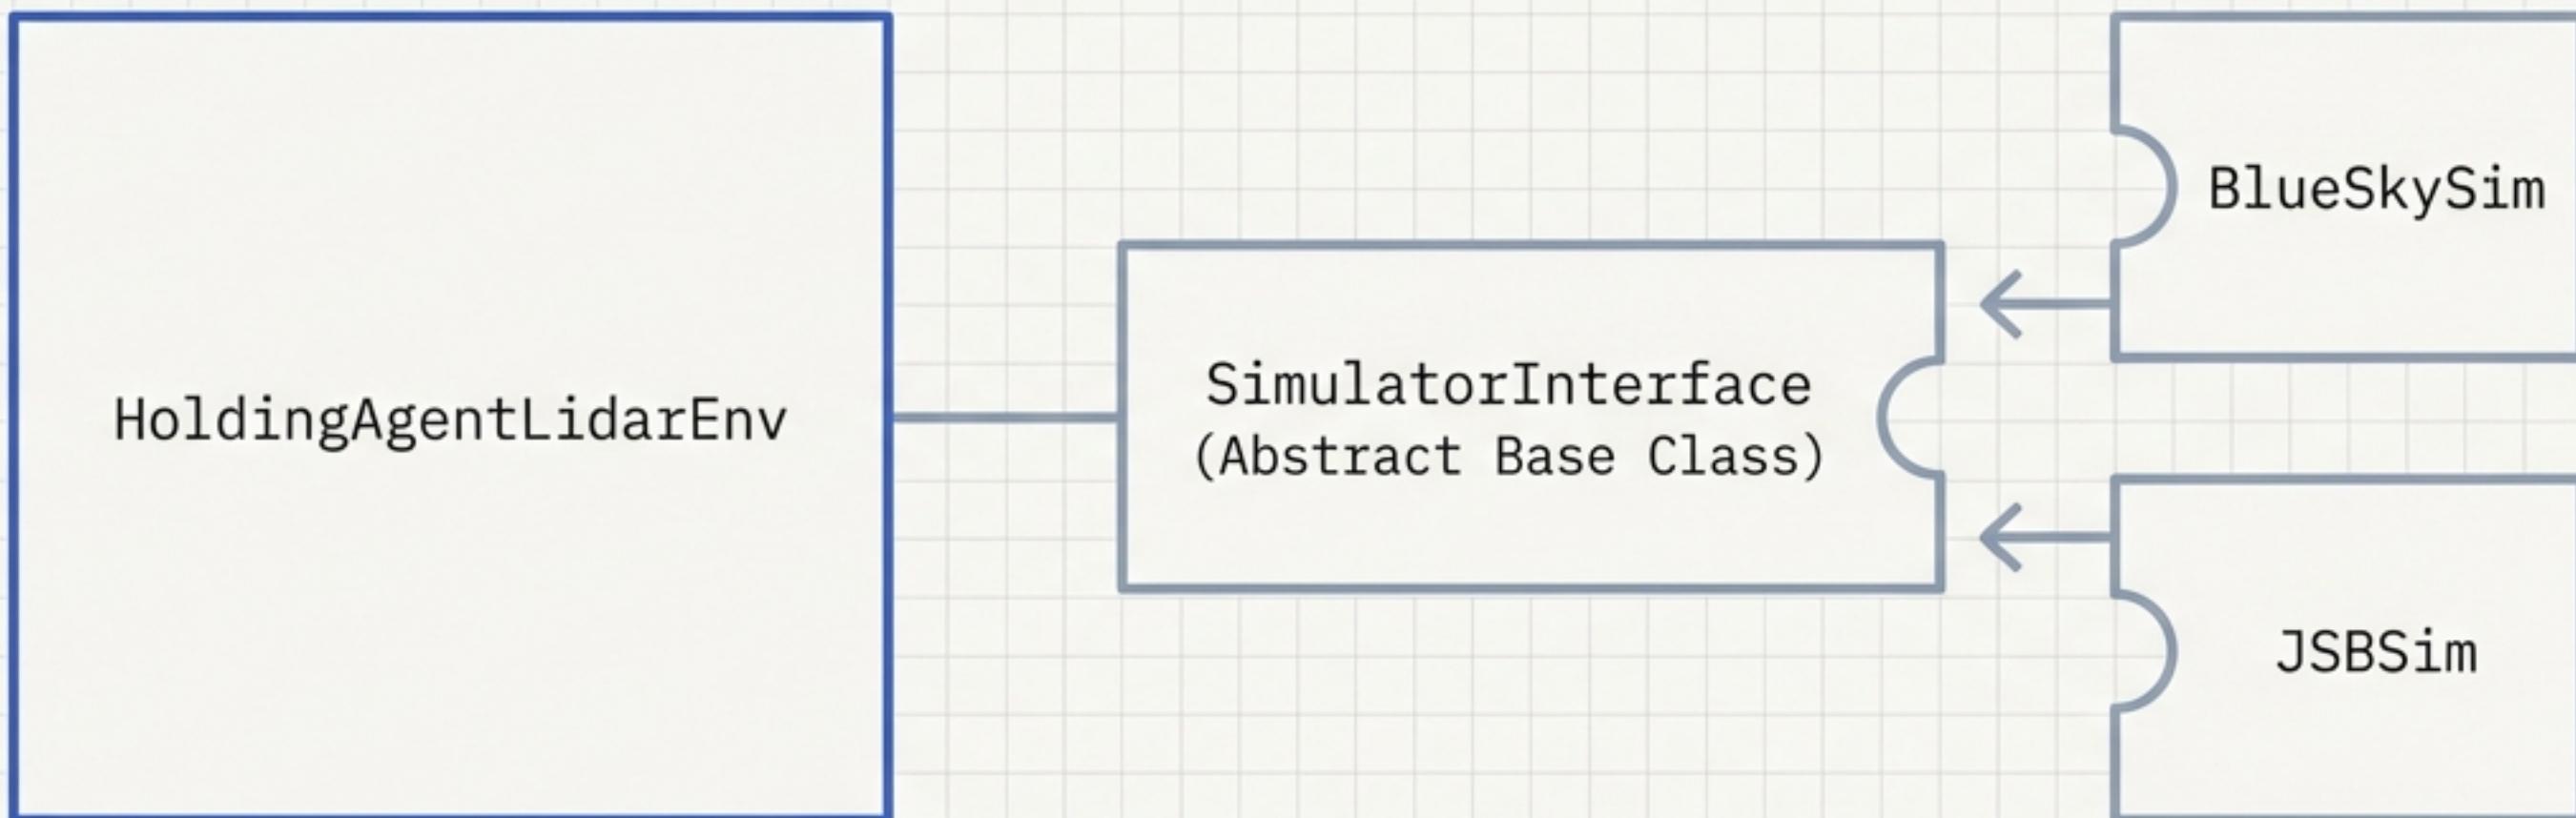
Isolate environment and simulator logic. Changes in one do not break the other, making the system easier to understand and evolve.



## REUSABILITY

Simulator interfaces and implementations can be shared across multiple projects and different RL environments.

# THE SYSTEM BLUEPRINT RELIES ON A FORMAL INTERFACE CONTRACT



The environment communicates through a `SimulatorInterface` contract, not directly with any simulator. Concrete simulator implementations act as plugins.

## THE `SIMULATORINTERFACE` DEFINES THE UNBREAKABLE CONTRACT

Any compliant simulator must implement the following abstract methods, guaranteeing predictable behaviour for the environment.

```
# The SimulatorInterface Contract
class SimulatorInterface(ABC):
    def reset(**kwargs)
    def step(dt)
    def get_aircraft_position(aircraft_id)
    def get_aircraft_heading(aircraft_id)
    def get_aircraft_speed(aircraft_id)
    def set_heading_change(aircraft_id, heading_change_deg)
    def close()
```

# SIMULATOR PORTABILITY IS A ONE-LINE CHANGE IN PRACTICE

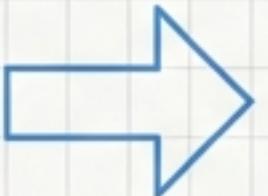
Because the environment is coded against the interface, switching the entire simulation backend is trivial.

BEFORE

```
# Using the BlueSky Simulator
from envs.simulators import BlueSkySim

sim = BlueSkySim()

env = HoldingAgentLidarEnv(simulator=sim)
```



AFTER

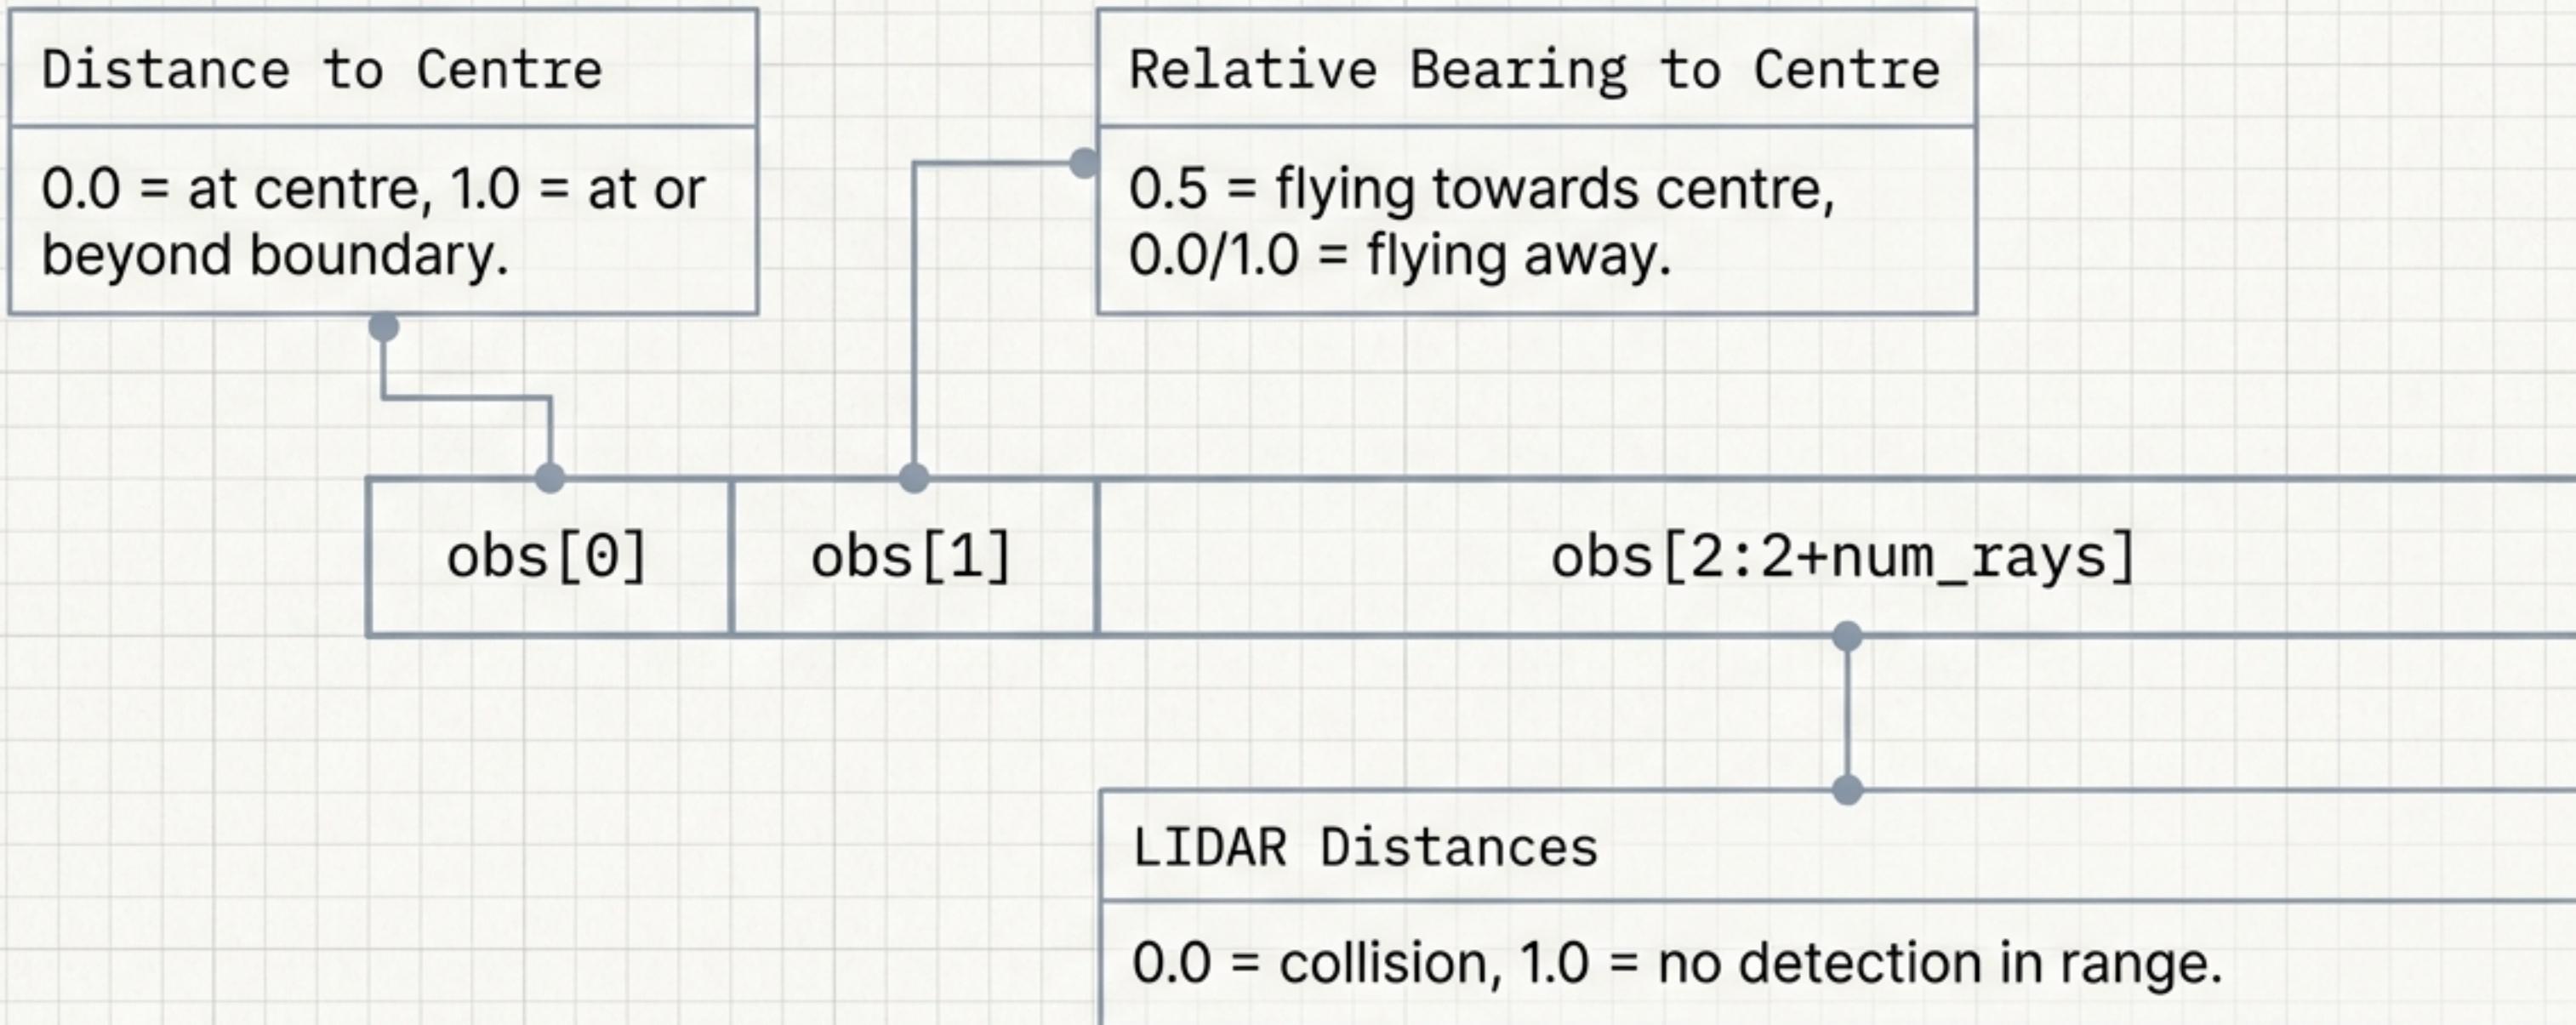
```
# Switching to the JSBSim Simulator
from envs.simulators import JSBSim

sim = JSBSim()

env = HoldingAgentLidarEnv(simulator=sim)
```

# THE AGENT PERCEIVES ITS WORLD THROUGH A COMPACT STATE VECTOR

The observation space is a continuous vector where each component is normalised to a [0.0, 1.0] range.



# THE AGENT ACTS BY COMMANDING A CONTINUOUS RATE OF TURN

## Type

Continuous, Box(1)

**Range:** [-1.0, 1.0]

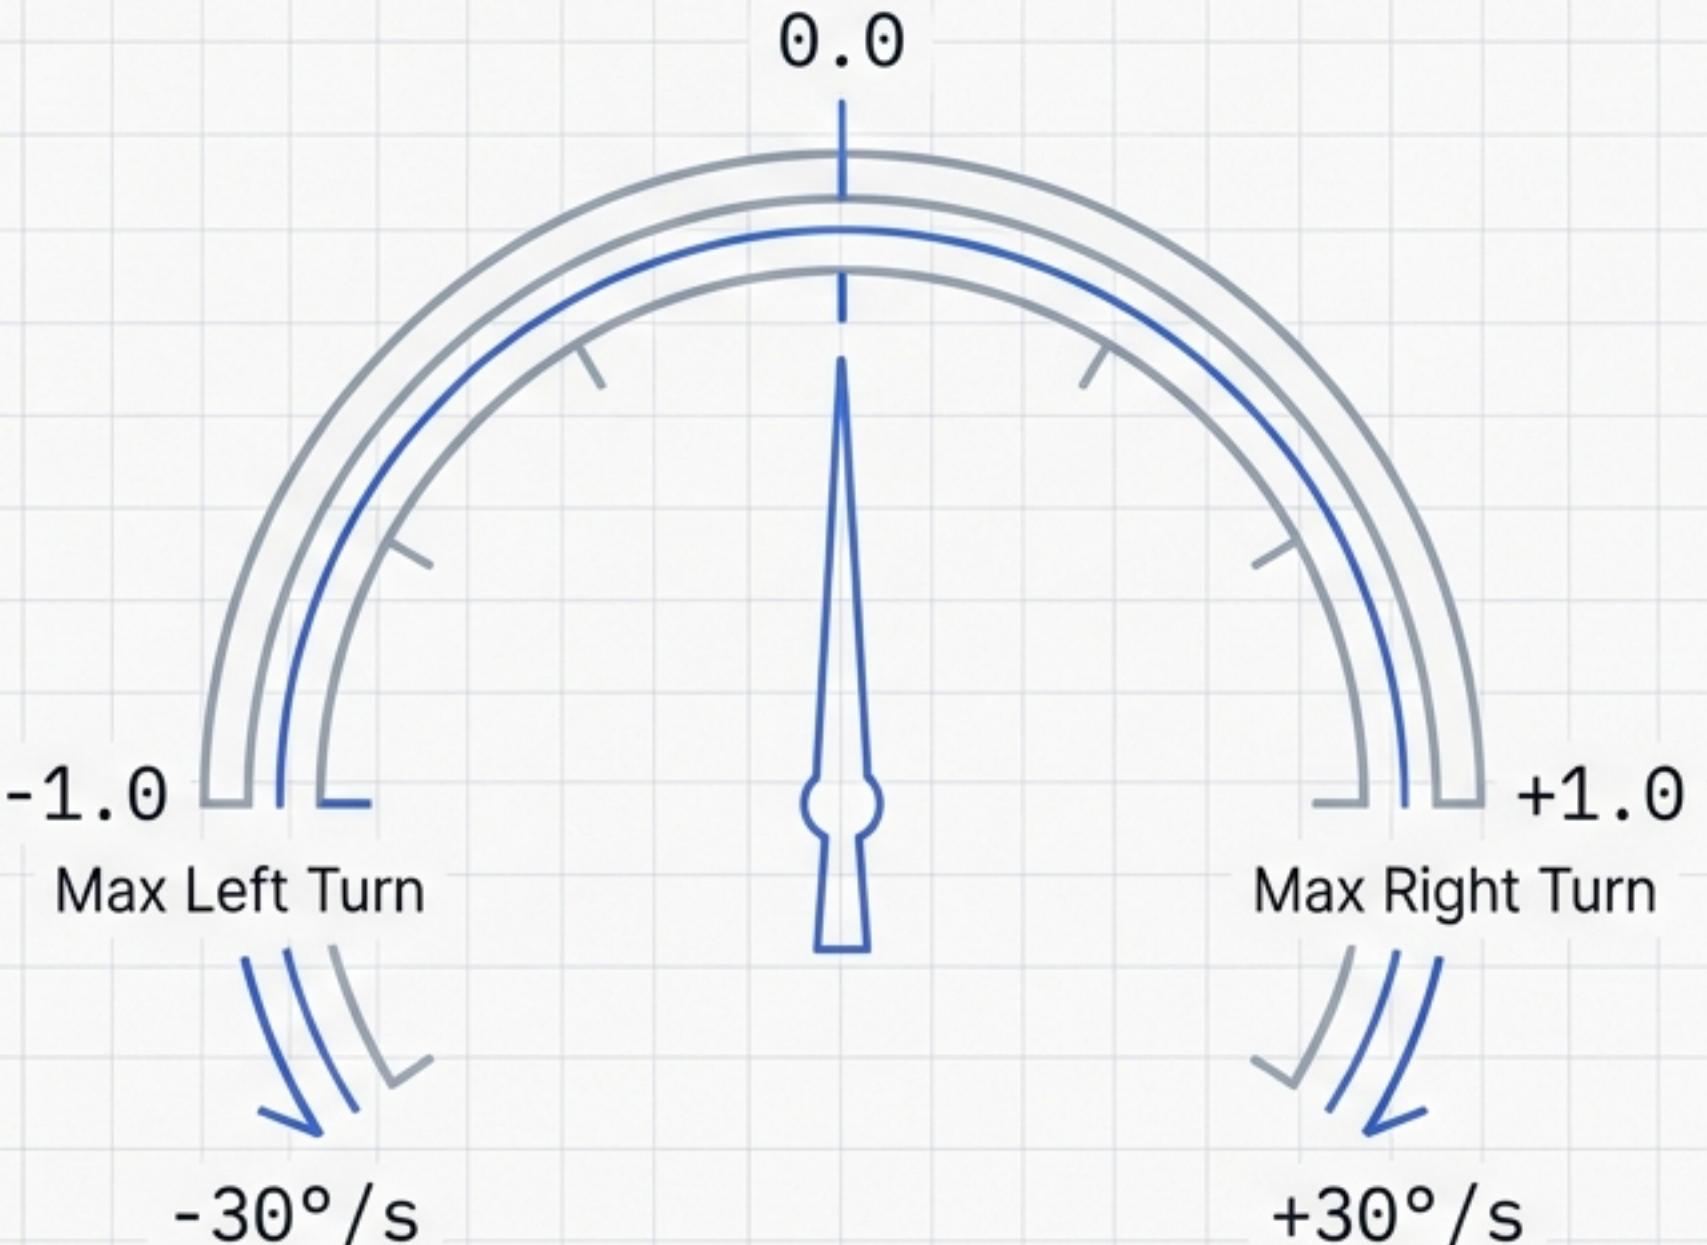
## Physical Meaning

The **action** value is scaled to a physical rate of heading change.

## Example

With a configured `turn_rate_deg_s` of 30:

- `action = 1.0` → Turn right at 30°/s
- `action = -1.0` → Turn left at 30°/s
- `action = 0.0` → Maintain current heading



# THE REWARD FUNCTION MOTIVATES SURVIVAL AND PUNISHES VIOLATIONS

Survival Reward	+0.1	Encourages longer episodes.
Boundary Penalty	-100	For leaving the containment zone.
Collision Penalty	-10	For entering an NFZ.
Proximity Penalty		A smooth negative reward that increases as the agent gets closer to a hazard, providing a useful learning gradient.

## Design Rationale

Violations are non-terminating. This is a deliberate choice to allow the agent to learn recovery behaviours, which is critical for robust, on-policy algorithms like PPO.

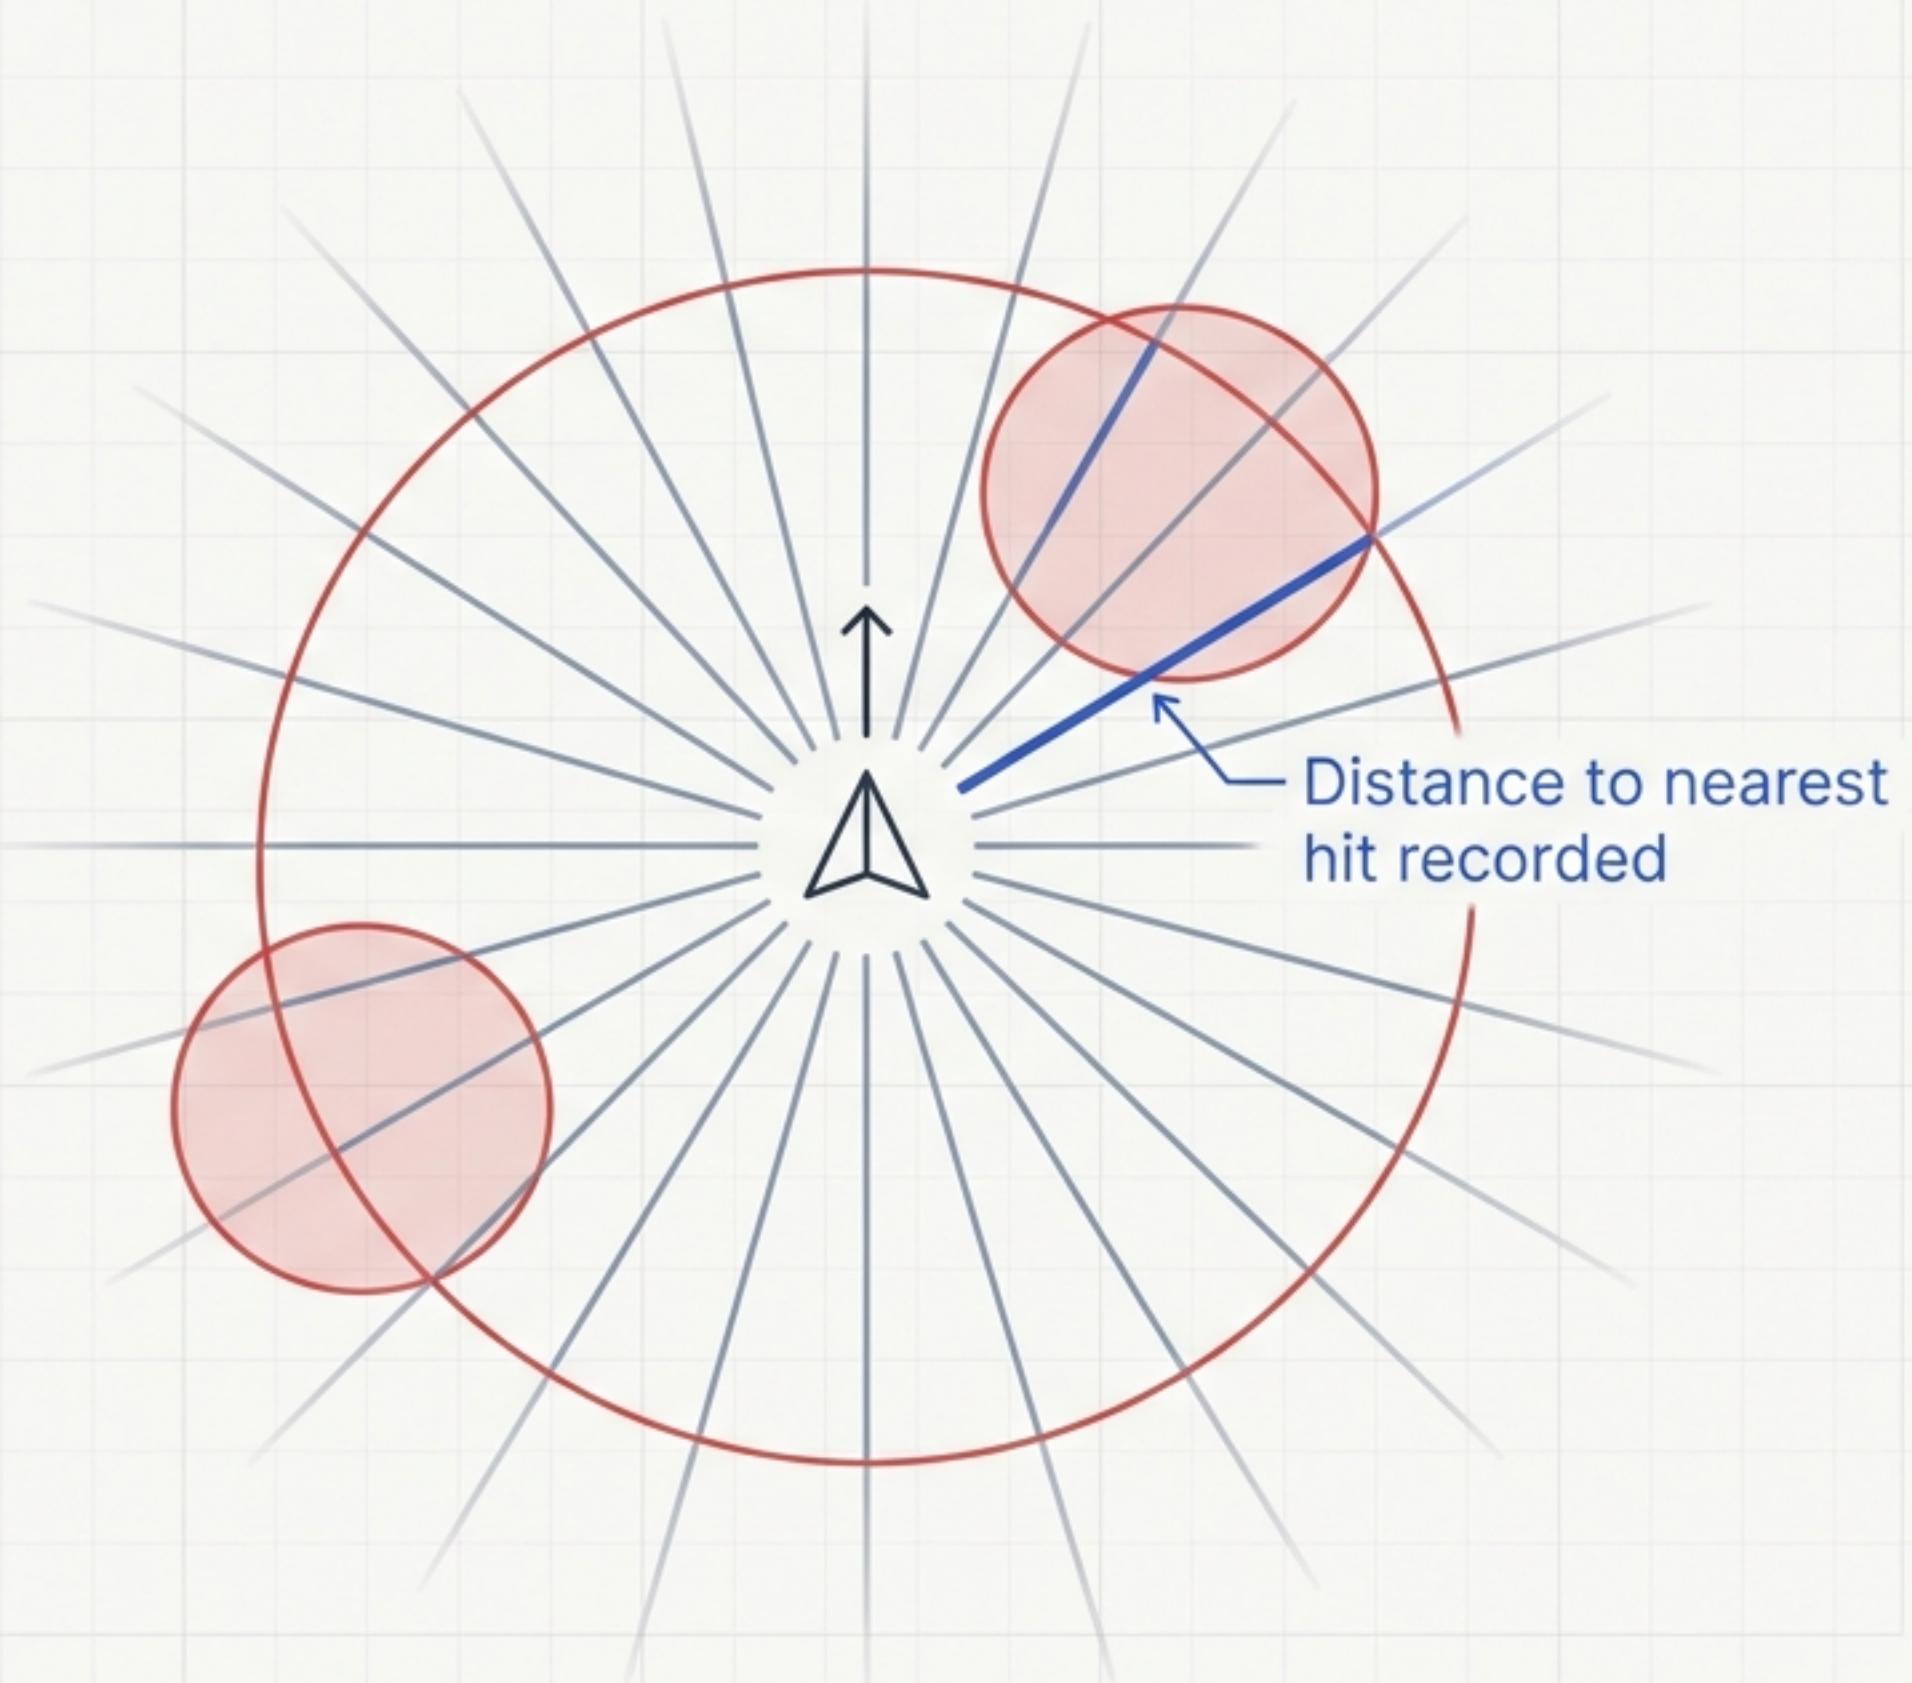
# A SIMULATED LIDAR PROVIDES 360° HAZARD PERCEPTION

The system casts a configurable number of rays radially around the aircraft.

Ray angles are relative to the aircraft's current heading (body frame).

Detection is based on efficient ray-circle intersection tests against all NFZs.

The system records only the distance to the *nearest* hit for each individual ray.



# COORDINATE SYSTEMS ARE HANDLED WITH A PRECISION-PERFORMANCE TRADE-OFF

The simulator provides **Latitude/Longitude**, but the environment's physics and LIDAR require a **local XY plane in metres**.

## Option 1: Pyproj (High Precision)



**Method:** Azimuthal Equidistant projection.

**Benefit:** Preserves true distances and bearings from the centre point. Geodetically accurate for any location or scale.

**Use Case:** Recommended for large areas (>50-100km radius).

## Option 2: Equirectangular (High Performance)



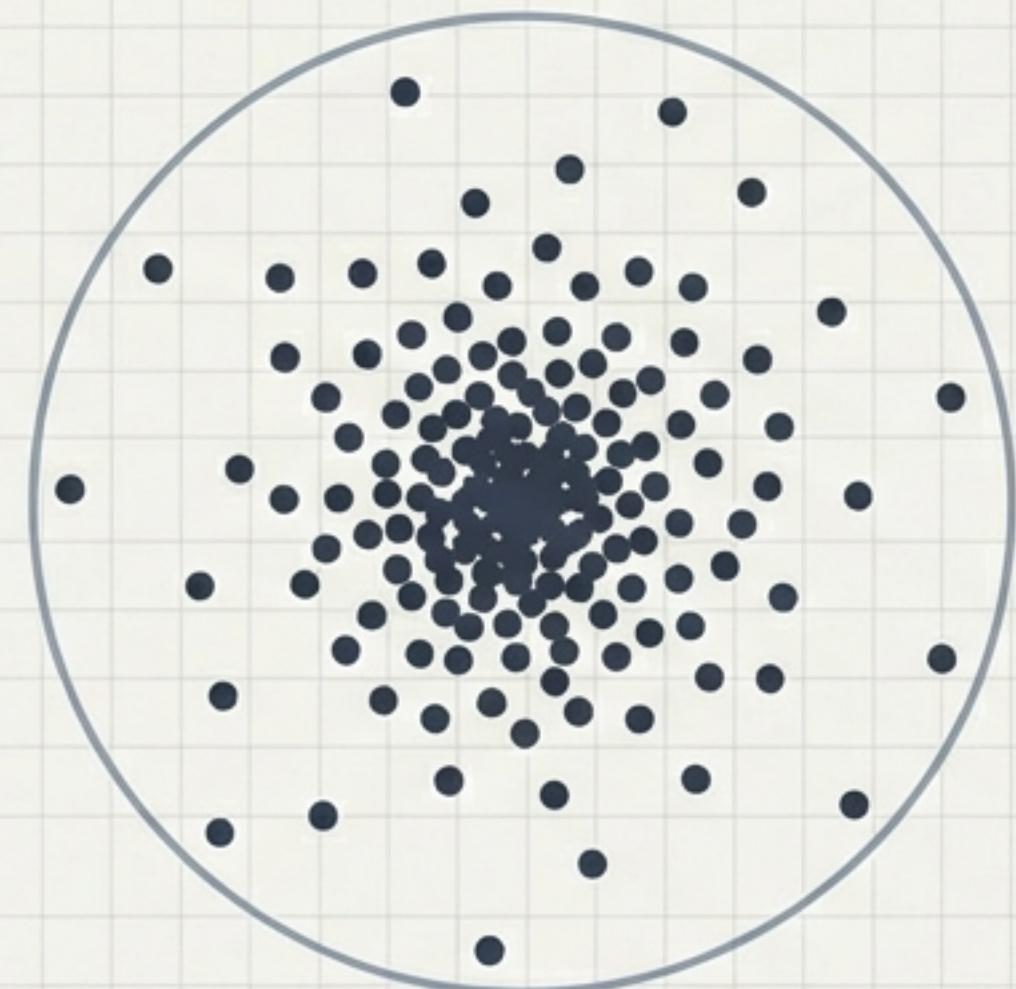
**Method:** A fast, linear approximation.

**Benefit:** No external dependencies, significantly faster computation.

**Use Case:** Excellent accuracy (<1% error) for areas up to ~100km at mid-latitudes (<60°).

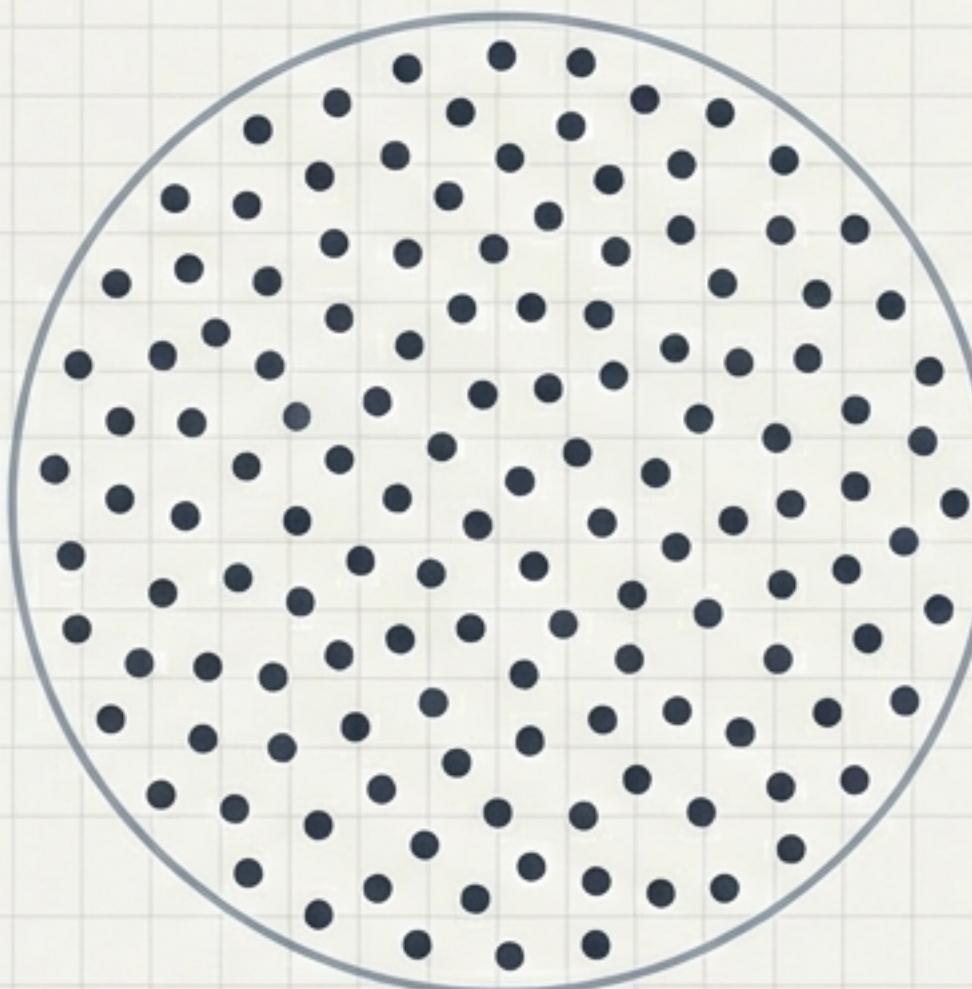
# Episode Randomisation Ensures Uniform Spatial Starting Positions

**Naive Sampling:** Biased to Centre



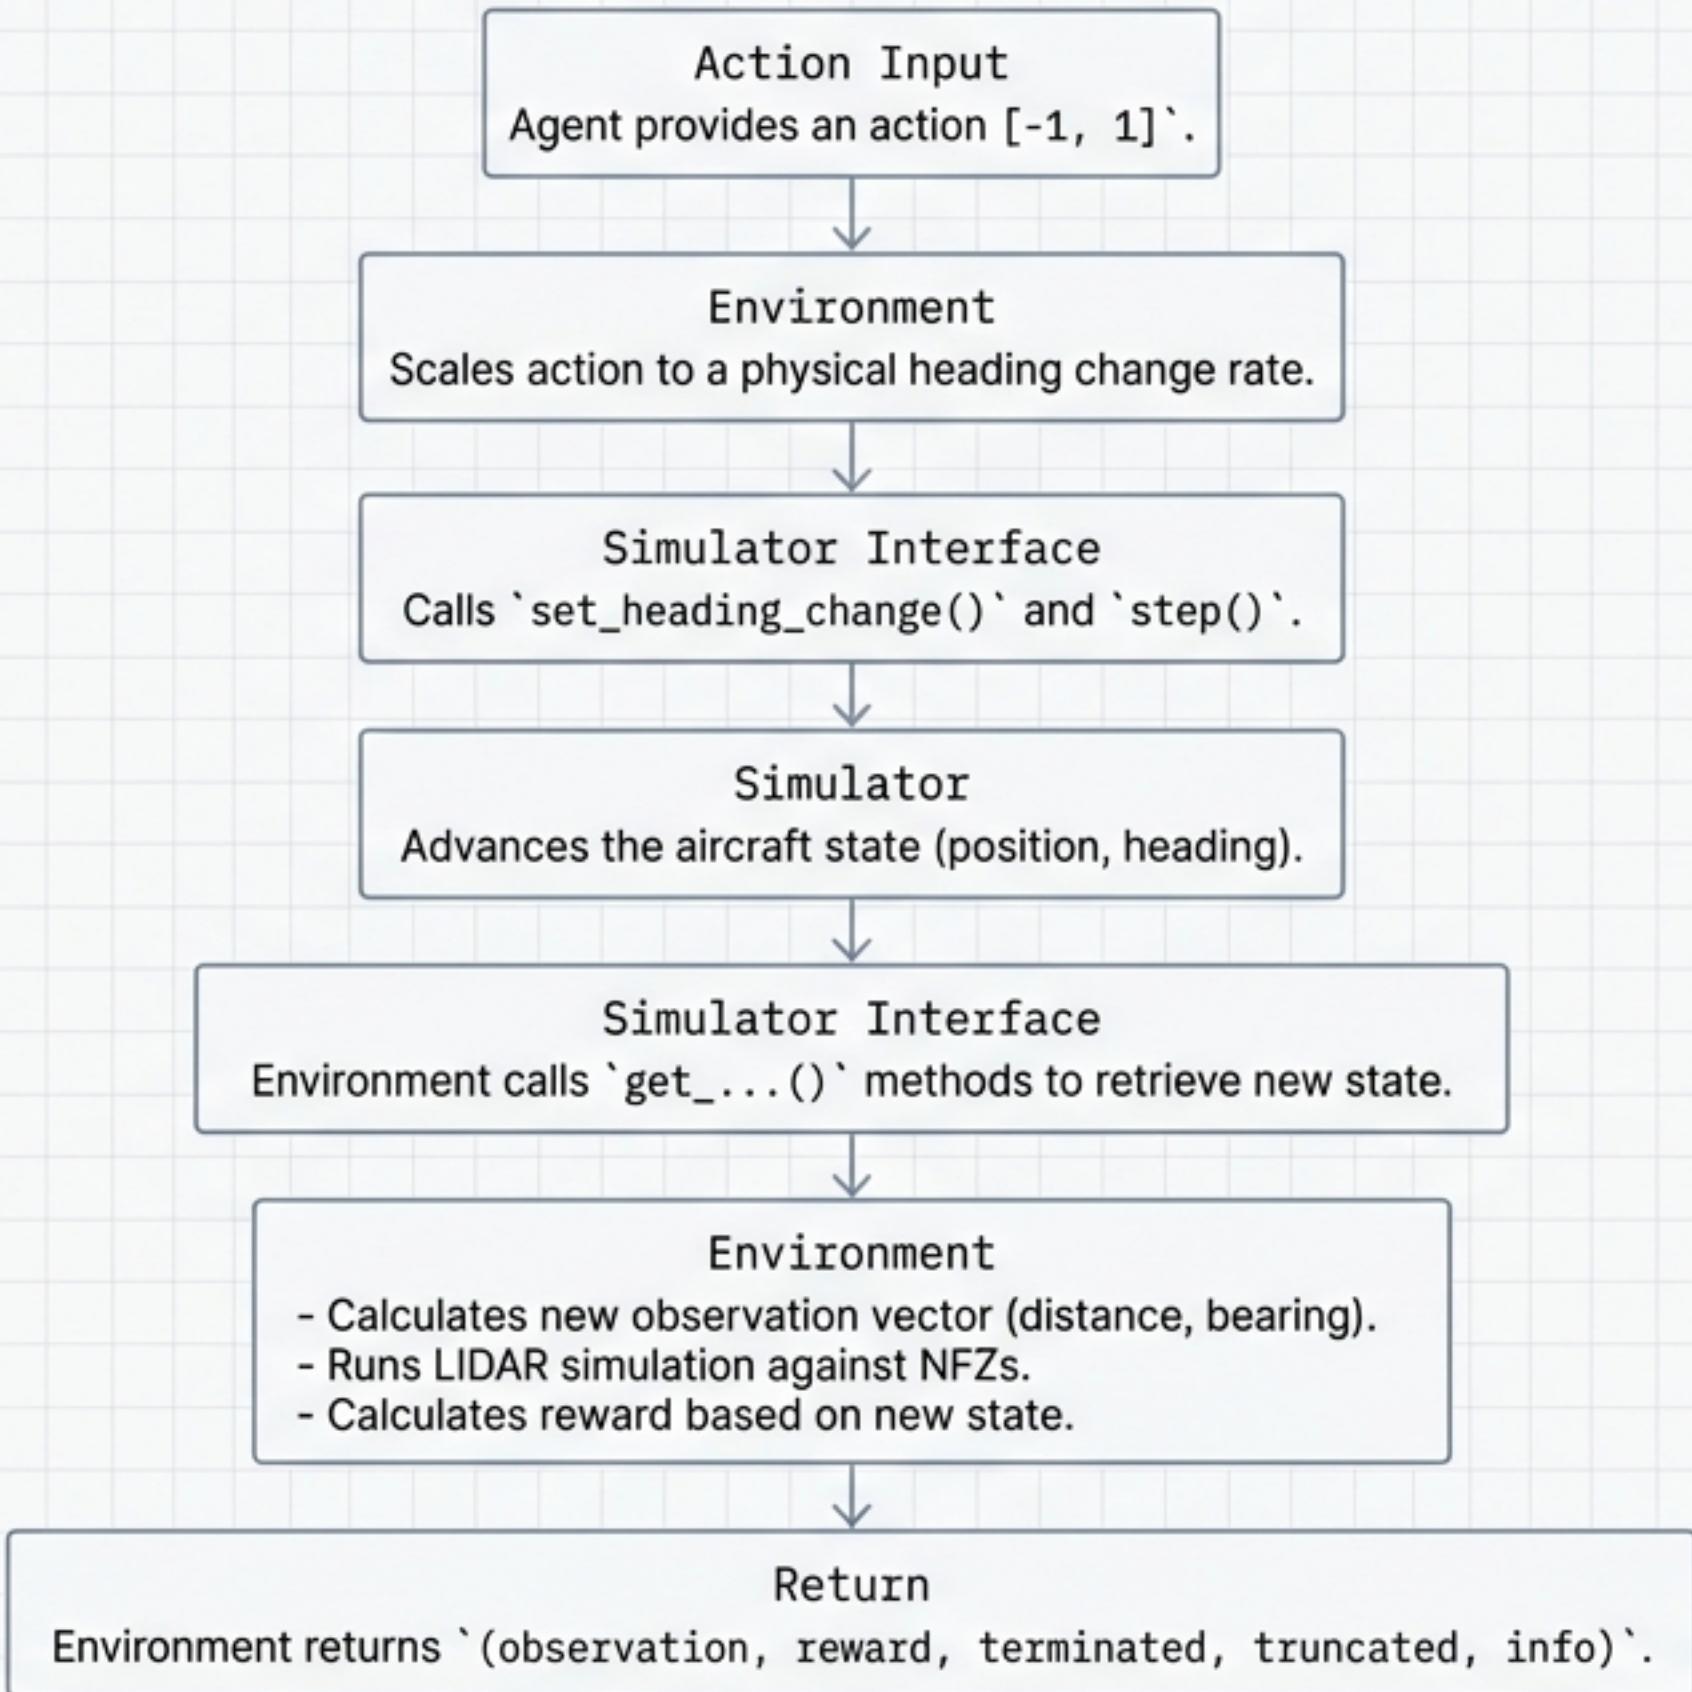
A naive randomisation of starting radius  $r = U(0, R)$  biases starting positions towards the centre of the holding zone.

**Correct Sampling:** Uniform Area Distribution



The radius is sampled using  $r = R * \sqrt{U(0, 1)}$ . This ensures that the agent's starting position is uniformly distributed over the entire circular area.

# The Step Execution Flow Orchestrates a Clear Sequence of Events



# Getting Started & Extending the Blueprint

## 1. Key Dependencies

- **Required:** gymnasium, numpy
- **Optional:** pyproj (for high-accuracy geodetics)

## 2. Actionable Next Steps

### To use with BlueSky:

1. Implement the methods in `envs/simulators/bluesky_sim.py`.
2. Connect to your BlueSky server instance.

### To add a new Simulator (e.g., JSBSim):

1. Create `envs/simulators/jsbsim_sim.py`.
2. Implement the `SimulatorInterface` contract.
3. Instantiate the environment with your new simulator class.

## 3.

